

# CONSERVATIVE GC

THE FINAL SEGFAULT

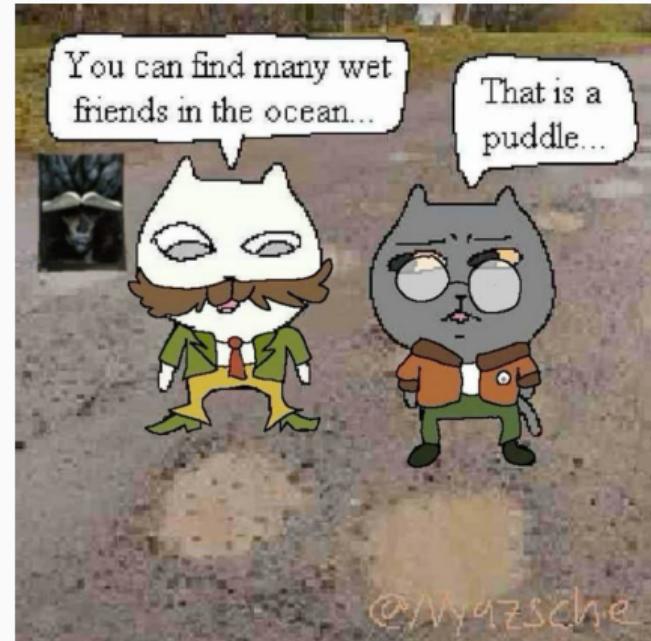
## GC: what is needed?

- Allocator
- Mark
- Sweep

# Allocator

## The whole API:

```
std::map<void*, size_t> allocated;  
void* heap;  
static constexpr size_t MIN_SIZE =  
    sizeof(void*) * 2;  
  
[[nodiscard]] void* allocate(size_t size);  
void deallocate(void* p);
```



## Mark

Find all reachable objects.

Search begins with GC roots, i.e.

- Globals
- Registers
- Stack



— [Logic – Challenging 12] I know how to find the roots.



— [Logic – Challenging 12] I know how to find the roots.  
**LOGIC** [Challenging: Failure] — You didn't think that scanning the globals and the stack would be so easy, right?

## Globals and stack

Turns out, finding the location of static memory is quite hard. There's no easy way to do it, only parsing the current ELF file and asking all of the loaded dynamic libraries.



# The solution



**AUTHORITY** [Trivial: Success] — Who cares about globals after all?  
Check only the registers and the stack.

# The solution



**AUTHORITY** [Trivial: Success] — Who cares about globals after all?  
Check only the registers and the stack.

**LOGIC** [Easy: Success] — In GC constructor save the current rsp value  
as the stack base and then iterate over the whole stack, finding the  
roots.

## Mark phase

**Mark needs two crucial things: looking into the memory behind found pointers and registering them as reachable.**

## Mark phase

Mark needs two crucial things: looking into the memory behind found pointers and registering them as reachable.



— How do I know the size of the object behind a pointer?

## Mark phase

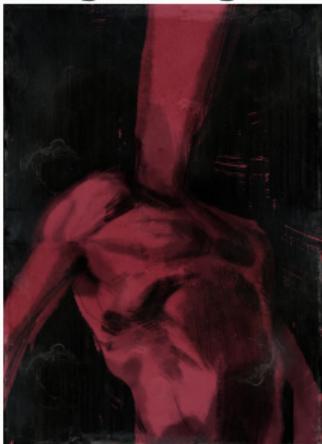
Mark needs two crucial things: looking into the memory behind found pointers and registering them as reachable.



— How do I know the size of the object behind a pointer?  
**PHYSICAL INSTRUMENT** [Medium: Success] — The %!@& do you need  
memory for? Look at those bytes. You \*must\* use them. You \*must\*  
rule them. Map is the way.

## Mark phase

Mark needs two crucial things: looking into the memory behind found pointers and registering them as reachable.



- How do I know the size of the object behind a pointer?  
**PHYSICAL INSTRUMENT** [Medium: Success] — The %!@& do you need memory for? Look at those bytes. You \*must\* use them. You \*must\* rule them. Map is the way.
- Sure thing. But how do we mark the objects that are reachable?

## Mark phase

Mark needs two crucial things: looking into the memory behind found pointers and registering them as reachable.



- How do I know the size of the object behind a pointer?  
**PHYSICAL INSTRUMENT** [Medium: Success] — The %!@& do you need memory for? Look at those bytes. You \*must\* use them. You \*must\* rule them. Map is the way.
- Sure thing. But how do we mark the objects that are reachable?  
**PHYSICAL INSTRUMENT** [Heroic: Success] — I'll say it again. The %!@& do you need memory for? Write a bitset for the whole heap.

## Sweep phase

Three lines:

```
for (auto&& [ptr, _) : allocator.allocated) {  
    if (!reachable[...]) {  
        allocator.deallocate(ptr);  
    }  
}
```


Just find all unreachable blocks and free them.

## Usage

We overloaded global operator 'new'.


## Usage

We overloaded global operator 'new'.  
That led to some funny consequences while  
testing: 'std::println' uses 'std::format',  
which, in turn, allocates a string, and its  
constructor is using 'new'.



## Usage

We overloaded global operator 'new'.  
That led to some funny cosequences while  
testing: 'std::println' uses 'std::format',  
which, in turn, allocates a string, and its  
constructor is using 'new'.  
But we used 'println' for logging in our  
allocator. We got to the point where doing  
literally anything with dynamic memory will  
cause a stack overflow.



Grand finale

# IT WORKS (kinda)



ABSOLUTE

## Harsh reality

We plugged in an AVL tree for testing and got all sorts of undefined behaviour.



## Harsh reality

We plugged in an AVL tree for testing and got all sorts of undefined behaviour. We fixed some UB. We saw some strange register shuffling. We saw the explosion of GoogleTest.



## Not so grand finale

We messed something up with scanning.  
Probably registers.



## Not so grand finale

We messed something up with scanning.  
Probably registers. Now we get a  
deterministic invalid behaviour.

